

**LLAMA 3  
IMPLEMENTED  
IN PURE  
NUMPY**



# OVERVIEW

As expected, the scale and performance is overwhelming. 24K GPUs, 15T training data, 10M instruction data, 1.3M GPU hours, it's all overwhelming. One interesting fact is that the model structure hasn't changed. Of course, Llama 3 have changed to using GQA, but this was already implemented in Llama 2 70B, so it's practically the same model structure.

We'll let it run for an accurate implementation, and we'll use only NumPy to make the model structure more intuitive to understand. We use the stories15M model that Andrej Karpathy trained while creating llama.2, by converting it to a NumPy compressed format using a converter. We will actually read in the model that Karpathy trained with the Llama 2 structure and implement it as executable code. One thing to note is that the stories15M model does not use GQA, so while we implement GQA in our code but not apply it to model behavior.

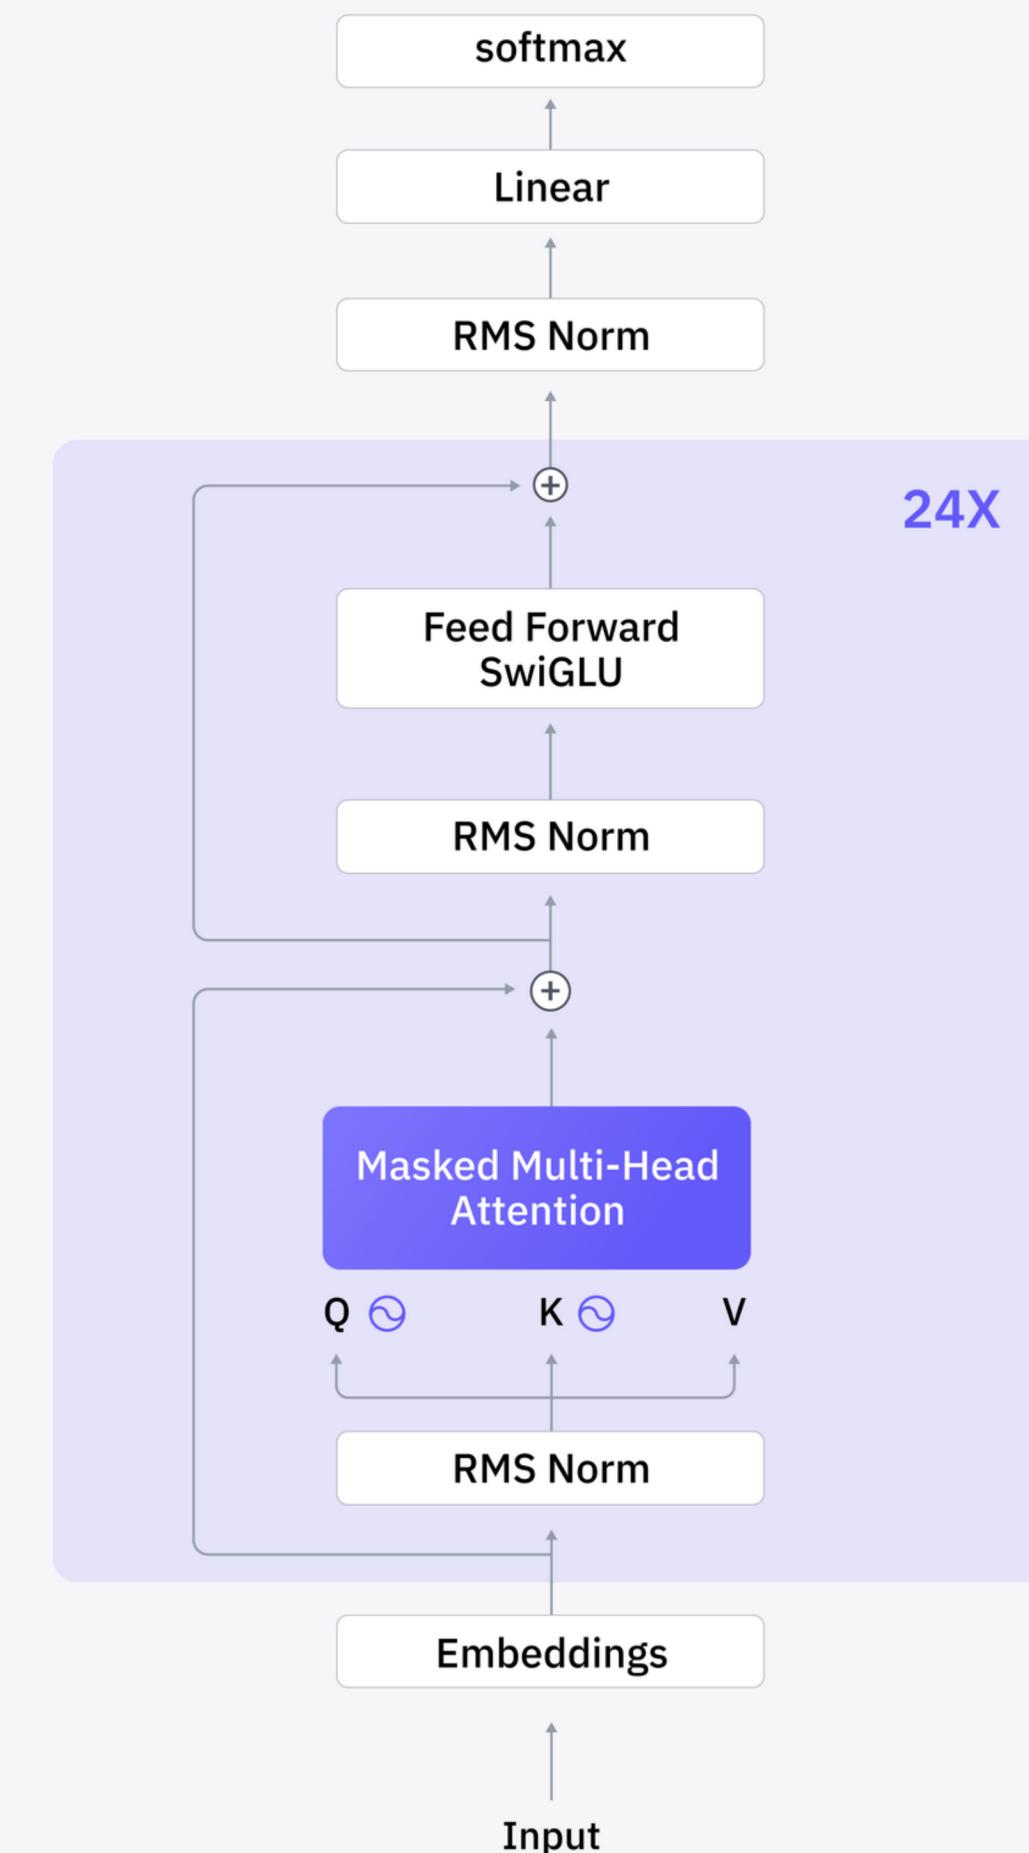
# STRUCTURE



The Model has the following parameters:

```
# Model params for ./stories15M.model.npz
dim: int                  = 288      # D
n_layers: int              = 6
n_heads: int               = 6        # QHN, HN, HD = 48
n_kv_heads: Optional[int]  = None     # KVHN = 6
vocab_size: int             = 32000    # VS
max_seq_len: int            = 256      # M
max_new_tokens: int         = 50
```

The designations D, HN, HD, VS, M etc. in the comments are used to manage the shape of each variable in code. Also note that unlike the 24x in the model illustration, stories15M model has 6 layers, so it iterates 6x.



# ROPE #1

The first step is to precompute cos and sin for RoPE embedding. These values are later used by `Q` and `K`. This calculation only needs to be done once for every request, so it can be cached. The size is `HD(48)//2`, which is an exponential multiple of `base(10000)`, so it can be a larger value, but the maximum value is never more than 1, so it is converted to a scaled value between `0 ~ 1`, and then again to a value between  $1 \sim \frac{1}{10000}$ .

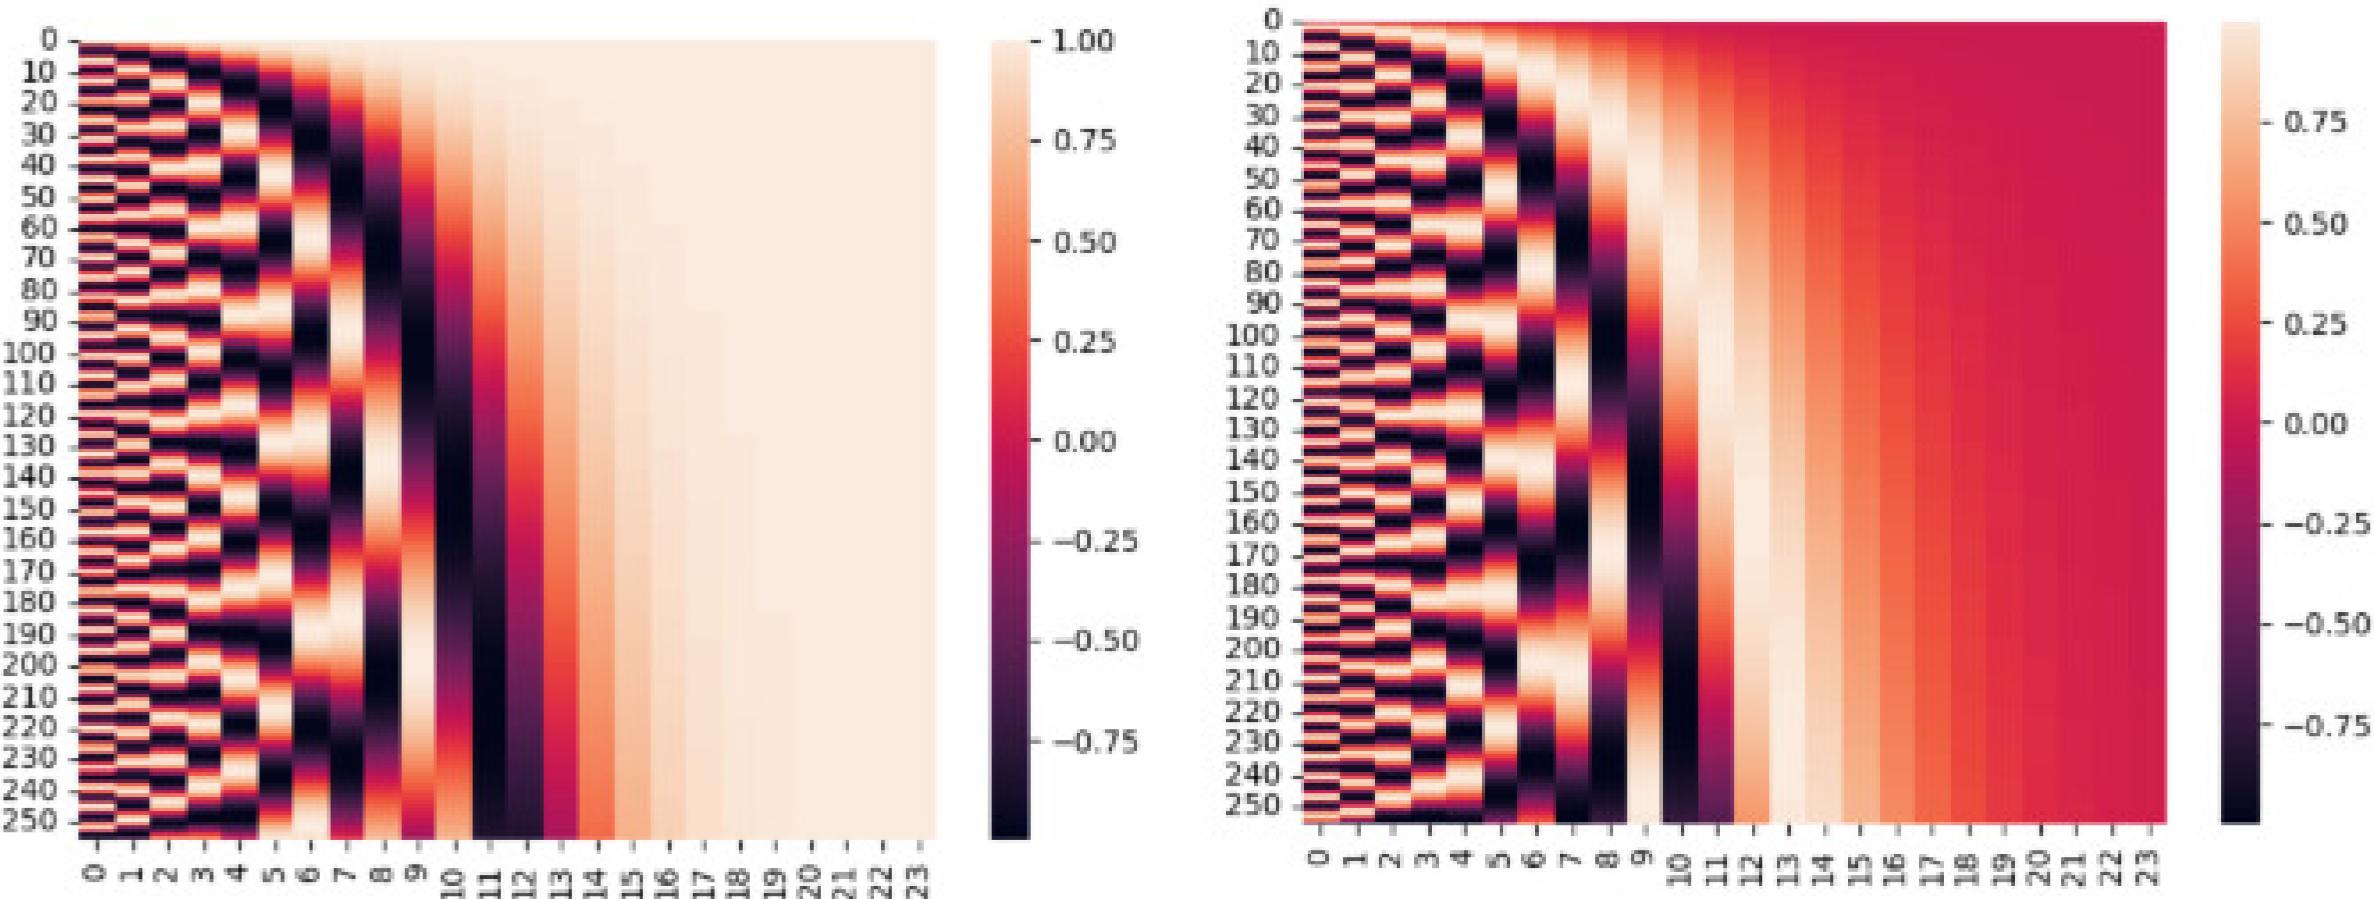
```
np.arange(0, 48, 2) # [24,]
1.0 / (base(10000) ** ([0, 2,          ..., 44,          46] / 48))
= 1.0 / (base(10000) ** [0, 0.04166667, ..., 0.9166667, 0.958333344])
= 1.0 /
[1, 1.4677993, ..., 4641.59, 6812.9194]
= [1, 0.68129206, ..., 0.00021544, 0.00014678]
```

The result of the calculation is `np.outer` multiplied by `max_seq_len(256)`, and then cos and sin are calculated.

```
# [256, ] x [24, ] = [256, 24]
freqs = np.outer([0 ~ 255], [1, 0.68129206, ..., 0.00021544, 0.00014678])
self.freqs_cos: Array["M, HD//2"] = np.cos(freqs)
self.freqs_sin: Array["M, HD//2"] = np.sin(freqs)
```

# ROPE #1

The heatmap of cos and sin looks like this:



The stories15M model is `max_seq_len(256)`, but I think it could scale up to 8K if we utilize all values up to horizontal axis 24.

# RMSNORM

RMSNorm normalizes activation values based on the Root Mean Square of the activation values, as opposed to using traditional Mini Batch or Layer statistics. This has the advantage of scaling activation consistently, regardless of Mini Batch size or Layer. Like other normalization techniques, it also has separate training parameters.

A well-known explanation of the success of LayerNorm is its re-centering and re-scaling invariance property. The former enables the model to be insensitive to shift noises on both inputs and weights, and the latter keeps the output representations intact when both inputs and weights are randomly scaled. In this paper, we hypothesize that the re-scaling invariance is the reason for success of LayerNorm, rather than re-centering invariance.

We propose RMSNorm which only focuses on re-scaling invariance and regularizes the summed inputs simply according to the root mean square (RMS) statistic:

$$\bar{a}_i = \frac{a_i}{\text{RMS}(a)} g_i, \quad \text{where } \text{RMS}(a) = \sqrt{\frac{1}{n} \sum_{i=1}^n a_i^2}. \quad (4)$$

Intuitively, RMSNorm simplifies LayerNorm by totally removing the mean statistic in Eq. (3) at the cost of sacrificing the invariance that mean normalization affords. When the mean of summed inputs is zero, RMSNorm is exactly equal to LayerNorm. Although RMSNorm does not re-center

# RMSNORM



The formula implementation is as follows:

```
z: Array["B, L or 1, 1"] = (x ** 2).mean(-1,  
keepdims=True) + self.eps  
z: Array["B, L or 1, D"] = x / np.sqrt(z)  
return z * self.weight
```

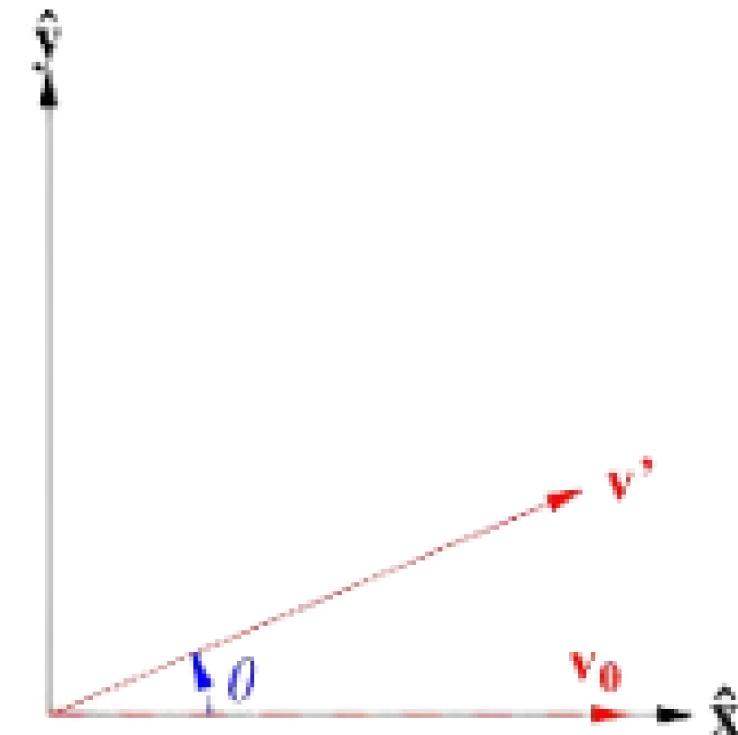
# QKV

The way to calculate QKV is to matmul one weight in GPT and then split it, but Llama have their own weights for QKV, so we need to matmul them separately. Then, for Multi-Head Attention, we reshape each one to separate them by Multi-Head.

```
● ● ●  
# QKV  
xq: Array["B, L or 1, D"] = x @ self.q_weight  
xk: Array["B, L or 1, D"] = x @ self.k_weight  
xv: Array["B, L or 1, D"] = x @ self.v_weight  
  
# ["B, L or 1, D"] -> ["B, L or 1, QHN or KVHN, HD"]  
xq: Array["B, L or 1, QHN, HD"] = xq.reshape(B, L,  
self.n_local_heads, self.head_dim)  
xk: Array["B, L or 1, KVHN, HD"] = xk.reshape(B, L,  
self.n_local_kv_heads, self.head_dim)  
xv: Array["B, L or 1, KVHN, HD"] = xv.reshape(B, L,  
self.n_local_kv_heads, self.head_dim)
```

# ROPE #2

Now it's time to actually apply the RoPE using the values we calculated earlier.



In  $\mathbb{R}^2$ , consider the matrix that rotates a given vector  $v_0$  by a counterclockwise **angle**  $\theta$  in a fixed coordinate system. Then

$$R_\theta = \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \quad (1)$$

so

$$v' = R_\theta v_0. \quad (2)$$

RoPE is a new type of position encoding technique that has the characteristics of both absolute and relative, and performs well because it has the characteristics of both. It only applies to Q and K, dividing each input by the sum of its parts, then multiplying by cos and sin, adding and subtracting the results, and returning the sum back to reshape.

# ROPE #2

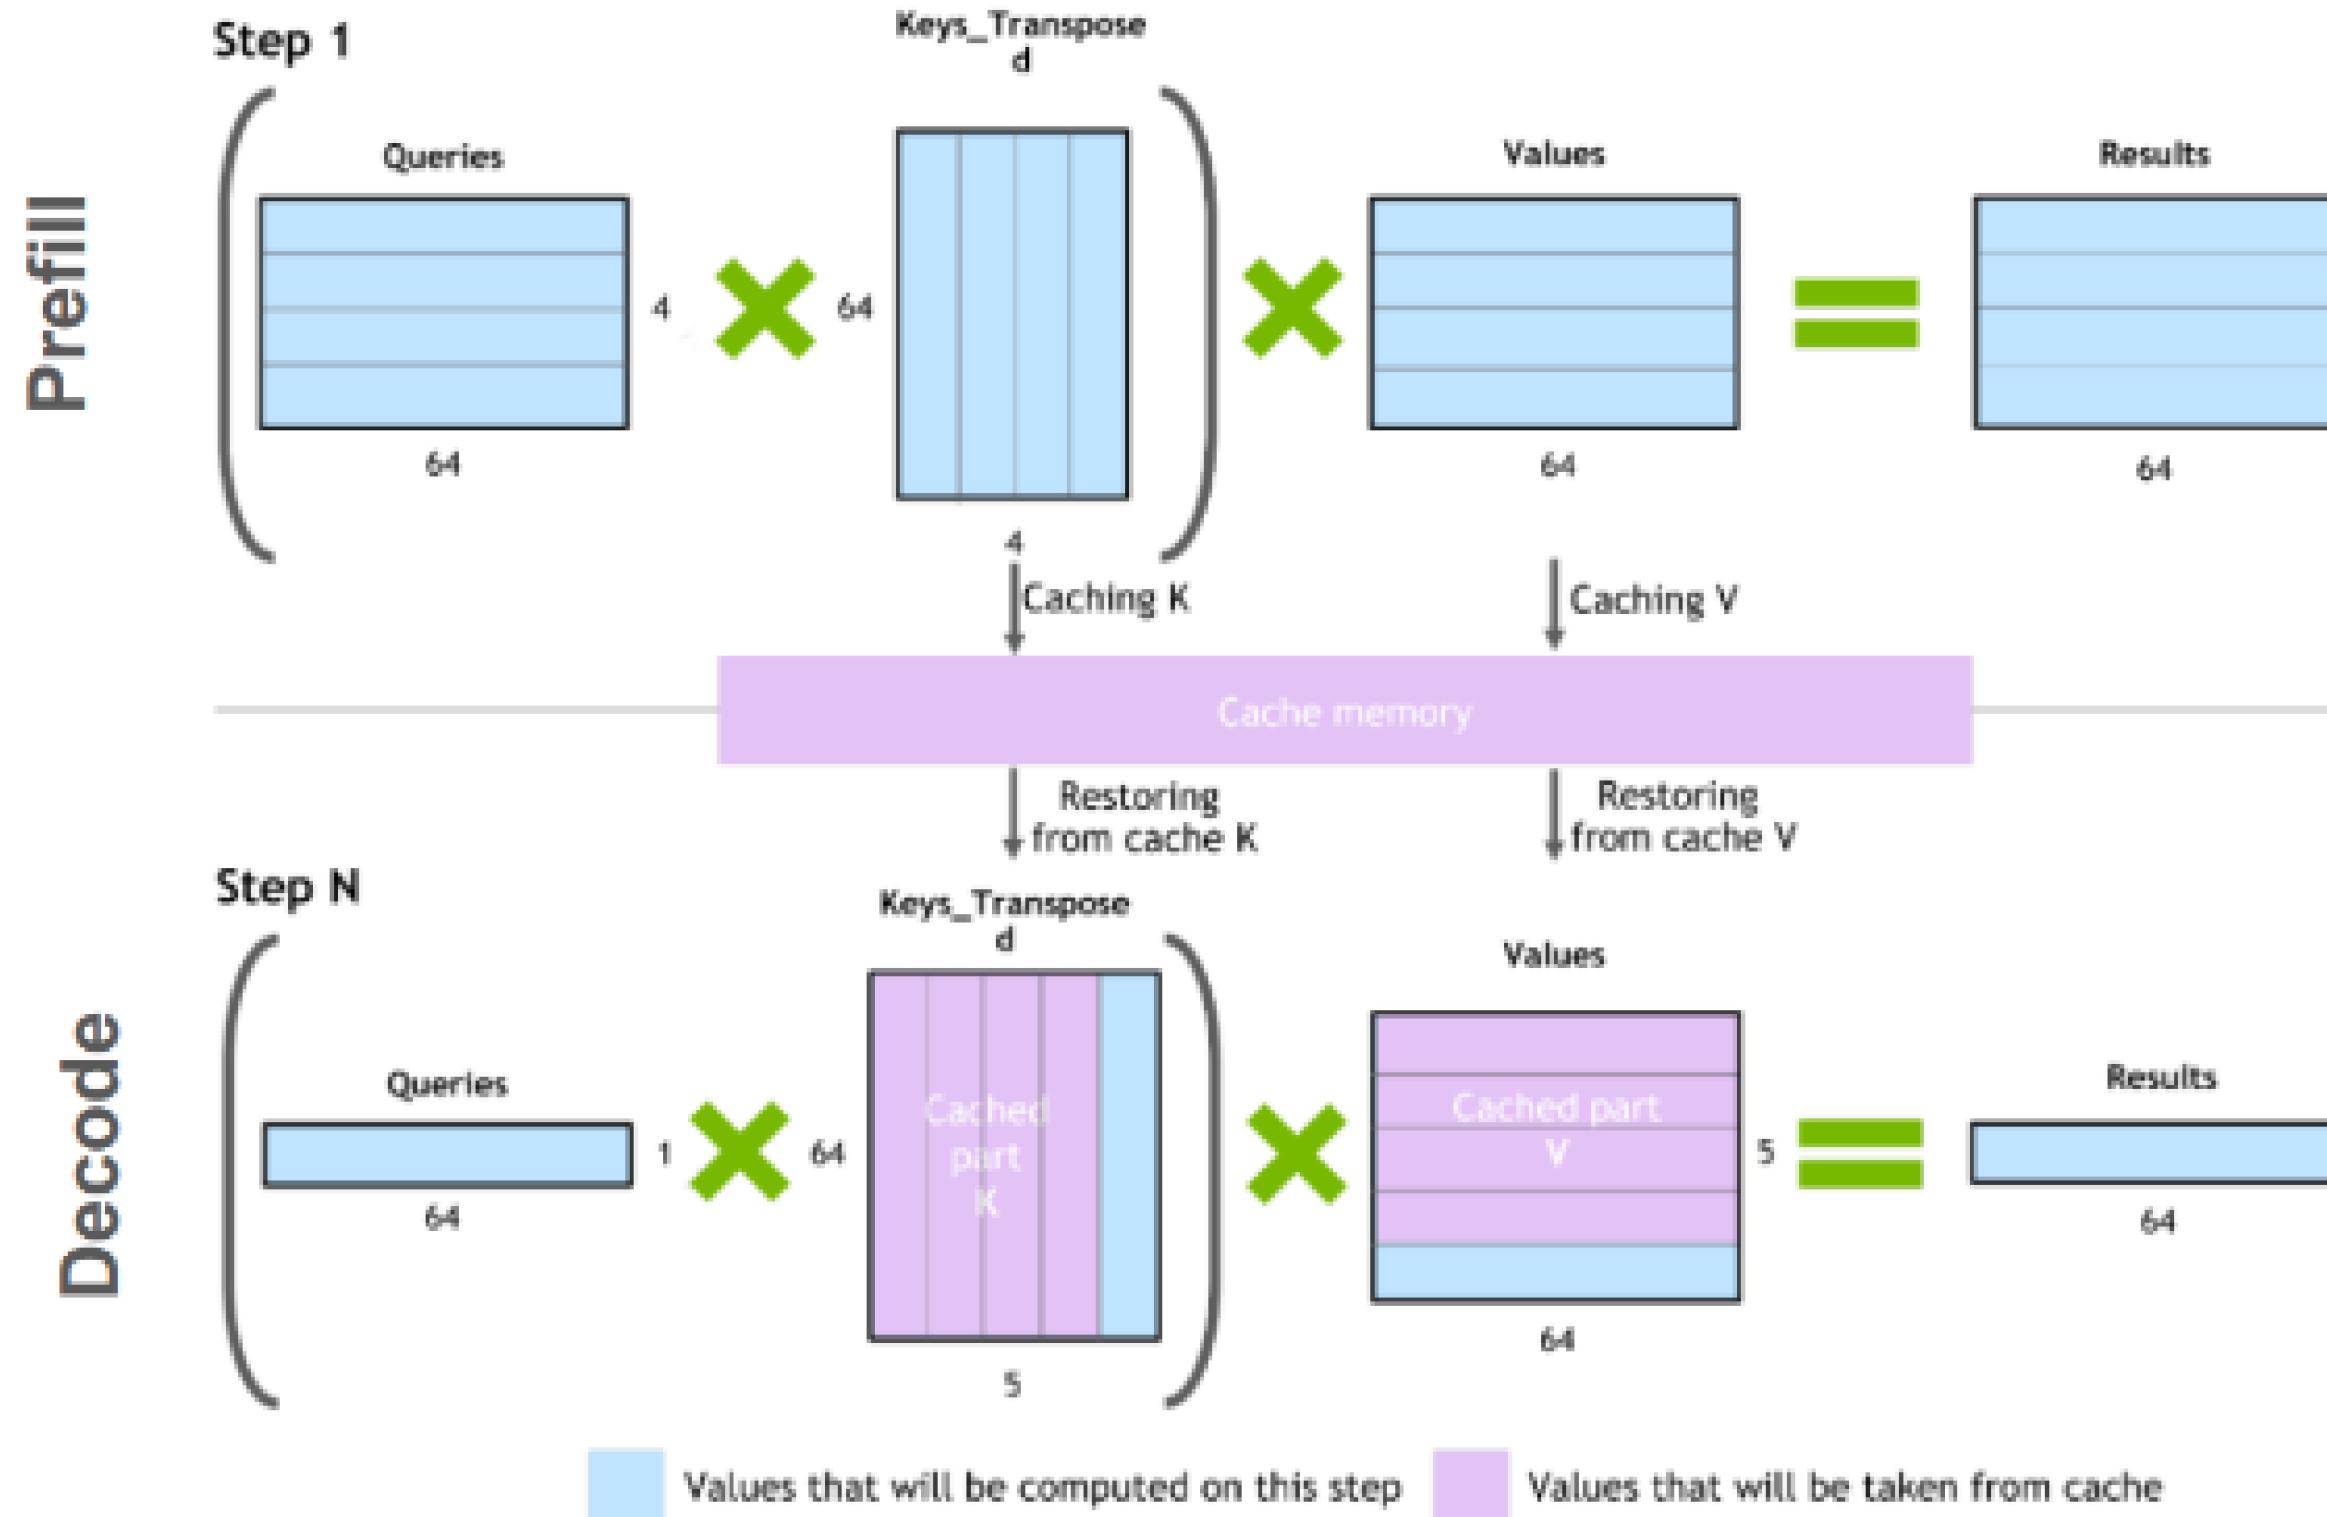


```
xq_out_r: Array["B, L or 1, QHN, HD//2"] = xq_r * freqs_cos - xq_i * freqs_sin  
xq_out_i: Array["B, L or 1, QHN, HD//2"] = xq_r * freqs_sin + xq_i * freqs_cos  
xk_out_r: Array["B, L or 1, KVHN, HD//2"] = xk_r * freqs_cos - xk_i * freqs_sin  
xk_out_i: Array["B, L or 1, KVHN, HD//2"] = xk_r * freqs_sin + xk_i * freqs_cos  
  
xq_out: Array["B, L or 1, QHN, HD//2, 2"] = np.stack([xq_out_r, xq_out_i], axis=-1)  
xk_out: Array["B, L or 1, KVHN, HD//2, 2"] = np.stack([xk_out_r, xk_out_i], axis=-1)  
xq_out: Array["B, L or 1, QHN, HD"] = xq_out.reshape(xq_out.shape[:-2] + (-1,))  
xk_out: Array["B, L or 1, KVHN, HD"] = xk_out.reshape(xk_out.shape[:-2] + (-1,))
```

RoPE are applied after the Q and K have been multiplied by the weights in the attention mechanism, while in the vanilla transformer they're applied before.

# KV CACHE

$(Q * K^T) * V$  computation process with caching



Since the GPT-style generative model is Masked Attention, it is possible to KV Cache. Since the previous result will always be the same, regardless of what comes after it, since we are not allowed to see the next word, we can cache K and V, and Q only needs to compute the last value. The cache is held by `max_seq_len(256)`, so the result of the calculation is put in and then extracted back to the only current length.

# KV CACHE

```
● ● ●

# KV Cache
self.cache_k[:B, start_pos: start_pos + L] = xk
self.cache_v[:B, start_pos: start_pos + L] = xv
ks: Array["B, L, KVHN, HD"] = self.cache_k[:B, : start_pos + L]
vs: Array["B, L, KVHN, HD"] = self.cache_v[:B, : start_pos + L]
# (1, 256, 6, 48) -> (1, 5, 6, 48)

# GQA
xk: Array["B, L, HN, HD"] = repeat_kv(ks, self.n_rep)
xv: Array["B, L, HN, HD"] = repeat_kv(vs, self.n_rep)

xq: Array["B, HN, L or 1, HD"] = xq.transpose(0, 2, 1, 3)
xk: Array["B, HN, L, HD"] = xk.transpose(0, 2, 1, 3)
xv: Array["B, HN, L, HD"] = xv.transpose(0, 2, 1, 3)
```

Here, we fetch the cache values and then transpose them back to reshape them, but this could be done more efficiently by skipping this step. For reference, the maximum size of the KV Cache is  **$1 \times 256 \times 6 \times 48 \times 2 \times 6 = 884$  K** on batch size 1. Since it is a 15M model, it takes up about 6% more memory.

# GQA(GROUPED-QUERY ATTENTION)

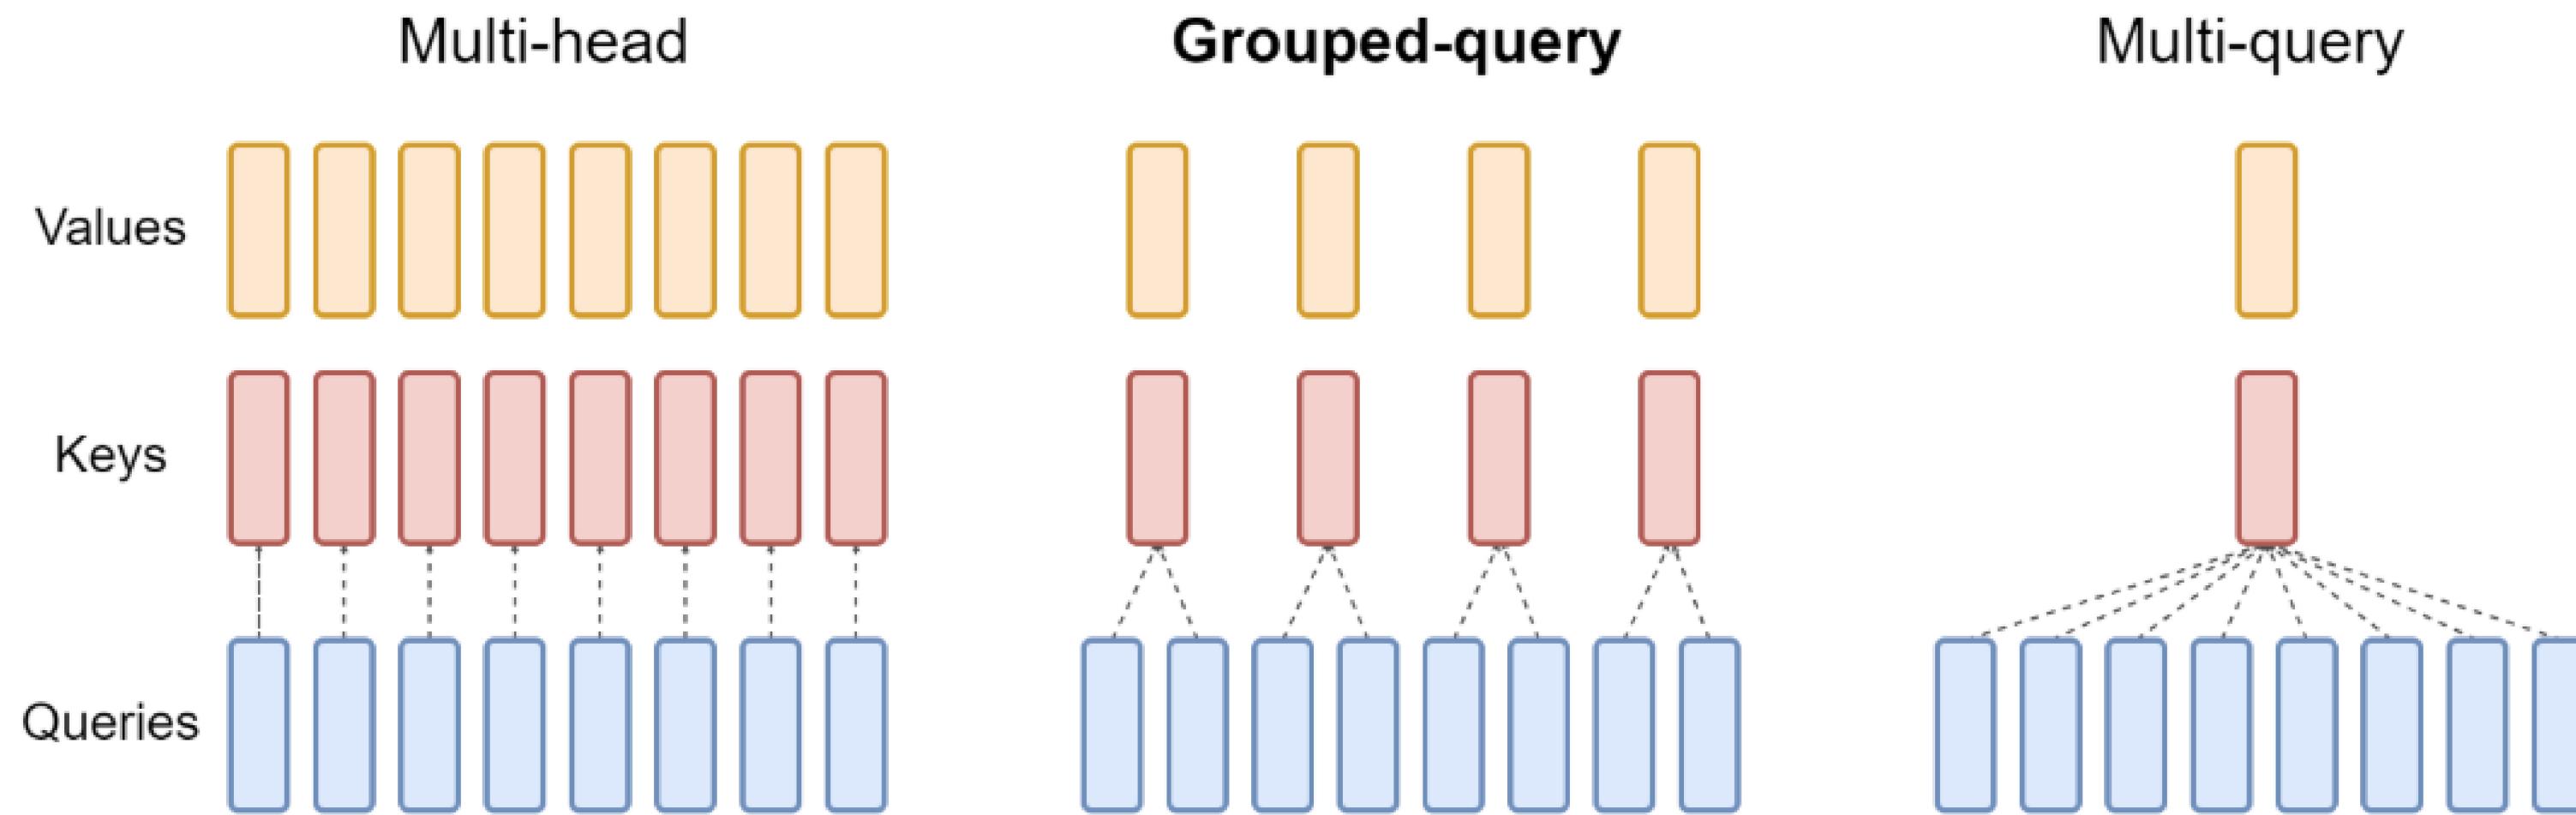


Figure 2: Overview of grouped-query method. Multi-head attention has  $H$  query, key, and value heads. Multi-query attention shares single key and value heads across all query heads. Grouped-query attention instead shares single key and value heads for each *group* of query heads, interpolating between multi-head and multi-query attention.

# GQA(GROUPED-QUERY ATTENTION)

```
● ● ●  
if n_rep == 1:  
    return x  
z: Array["B, L, QHN, HD"] = np.repeat(x, n_rep, axis=2)
```

MQA, which is a Multi-query, has the advantage of being compact and memory-saving compared to MHA, which is a Multi-head, but it suffers from poor performance and unstable learning. Therefore, Grouped-query, GQA, was introduced in Llama 2. In Llama 2, GQA was only applied to 70B, but from Llama 3, GQA was applied to all models above 8B. Since we are using a model that was trained without GQA, we do not use GQA, but we have implemented it in the code. We have implemented it by simply copying it by a multiple, and it can be improved by referencing the previous value for future optimization. We have avoided using GQA when n\_rep==1.

# SCALED DOT-PRODUCT ATTENTION

Attentions are calculated separately by Multi-Head.

$$\text{Attention}(Q, K, IN) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)IN$$



```
# Scaled Dot-Product Attention
# ["B, HN, L or 1, HD"] @ ["B, HN, HD, L"] -> ["B, HN, L or 1, L"]
attention: Array["B, HN, L or 1, L"] = xq @ xk.transpose(0, 1, 3, 2) / math.sqrt(self.head_dim)
# `mask` is used only once at the beginning.
if mask is not None:
    attention = attention + mask[None, None, :, :]
attention = softmax(attention)
output: Array["B, HN, L or 1, HD"] = attention @ xv
```

Masking is only done at the beginning and only the last Q needs to be processed afterward, so no masking is needed. The result can then be obtained with softmax and matmul. Finally, the result of the Multi-Head calculation is reshaped to full dimension to combine the heads and matmul once more.



```
# ["B, HN, L or 1, HD"] -> ["B, L or 1, D"]
output: Array["B, L or 1, D"] = output.transpose(0, 2, 1, 3).reshape(B, L, -1)
output: Array["B, L or 1, D"] = output @ self.o_weight
```

Computing the entire QKV at once is only done in the Prefill Phase. At this time, TTFT (Time To First Token) is called Prefill Latency, and only ‘vector @ matrix’ operations need to be performed from the Decode Phase onward. Flash Attention is also effective only when reducing the Prefill Latency during inference, and it performs somewhat well when the input is long.

# FEED FORWARD

In Llama model, Feed Forward uses 3 linear with matmul only and no bias, so unlike GPT, it is not a complete fully-connected layer. We create a swish value from the silu result, multiply it with  $x_V$  up-scaled from D to FD, and down-scale it again. Here, the size of FD is  $FD = 2 * 4 * D / 3$ , which is D(288), so FD(768).



```
swish: Array["B, L or 1, FD"] = silu(x @ self.gate_weight)
x_V: Array["B, L or 1, FD"] = x @ self.up_weight
x: Array["B, L or 1, FD"] = swish * x_V
x: Array["B, L or 1, D"] = x @ self.down_weight
```

# SwiGLU

In the paper, the SwiGLU formula looks like this:

**Transformer** ("Attention is all you need")

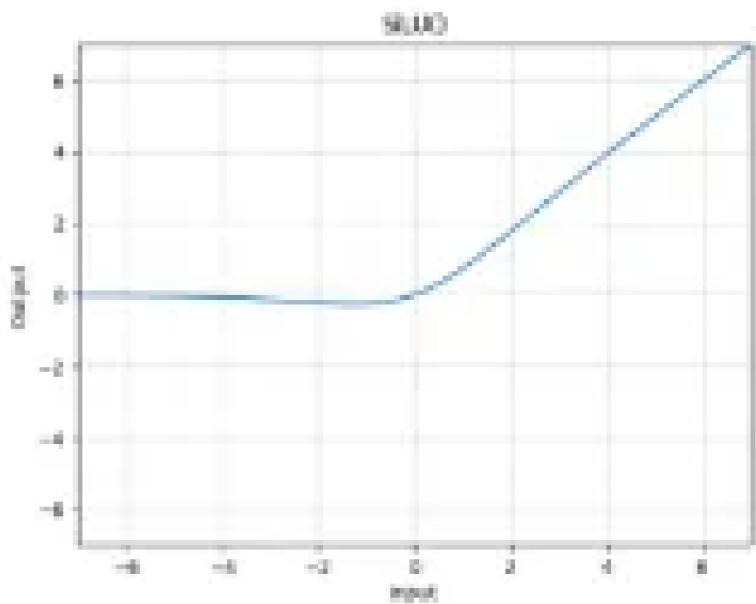
$$\text{FFN}(x) = \max(0, xW_1 + b_1)W_2 + b_2$$

**LLaMA**

$$\text{FFN}_{\text{SwiGLU}}(x, W, V, W_2) = (\text{Swish}_1(xW) \otimes xV)W_2$$

We use the swish function with  $\beta = 1$ . In this case it's called the Sigmoid Linear Unit (SiLU) function.

$$\text{swish}(x) = x \text{ sigmoid}(\beta x) = \frac{x}{1 + e^{-\beta x}}$$



Multiplying  $x_V$  with swish and matmul it with  $W_2$  is called SwiGLU. This unique combination of multiple feed forwards layers increases the performance of the model. Let  $x$  be a real number between approximately  $-14 \quad 11$ , which is the input to the silu function. The silu implementation is as follows:

```
x * (1 / (1 + np.exp(-x)))
```

# LINEAR

After passing through all the transformer blocks, the final output is only the last logit computed by matmul to speed things up. The transformer block always outputs ["1, D"] as the result after the Prefill Phase.

```
● ● ●  
# ["B, 1, VS"] = ["B, 1(L), D"] @ ["D,  
VS"]  
logit: Array["B, 1, VS"] = h[:, [-1], :]  
@ self.lm_head_weight
```

# GENERATION

Now, we generate tokens one after the other using the extracted logit. For simplicity, we've omitted sampling from the generation process and only output the Greedy result.

```
● ● ●  
  
for i, curr_pos in enumerate(range(L, max_new_tokens)):  
    if i == 0: # Prefill Phase  
        inputs = input_ids  
        pos = 0  
    else:       # Decode Phase  
        inputs = next_id  
        pos = curr_pos  
    logits: Array["B, 1, VS"] = self(inputs, pos)  
    next_id = logits[:, -1, :].argmax(-1, keepdims=True)  
    yield next_id
```

- The first step is Prefill Phase, or sometimes called Summarization. It passes all input and starts at position 0. This is also where Flash Attention comes into play.
- From then on, it is the Decode Phase and thanks to the KV Cache, only the last token ID is passed to Q and the result is the also last logit. Here, we omit sampling and only extract the maximum value. If you want to add a sampling process, you can take softmax and implement top\_p and top\_k.
- You can now yield the token ID we generated as a result, decode it in the next step and print the output token to finalize the process.

# EXAMPLE

## You can run it like this:

```
● ● ●

$ python llama3.py "I have a dream"
"""

I have a dream. He dream of a big, beautiful garden full of flower
and tree. He dream of playing with hi friend and eating yummy snack.
One day, he wa walking in the garden when he saw

Token count: 50, elapsed: 1.53s, 33 tokens/s
"""

|
```

# likejazz/llama3.np

llama3.np is pure NumPy implementation for Llama 3 model.



1

Contributor

1

Issue

408

Stars

23

Forks



## likejazz/llama3.np: llama3.np is pure NumPy implementation for Llama 3 model.

llama3.np is pure NumPy implementation for Llama 3 model. -  
[likejazz/llama3.np](https://github.com/likejazz/llama3.np)

GitHub